

PostgreSQL 17 Database Administration: Mastering `max_connections` and Connection Management

J

[Jeyaram Ayyalusamy](#)

Following

14 min read

.

Jun 12, 2025

4

2

Press enter or click to view image in full size



PostgreSQL is one of the most advanced open-source relational database systems, but many DBAs overlook one simple parameter that can make or break performance: `max_connections`. Managing connections properly is

essential for scaling PostgreSQL safely, especially in PostgreSQL 17 where workloads are heavier and systems are often more demanding.

In this post, we'll deep-dive into how PostgreSQL handles connections, potential pitfalls, and practical steps you can take to optimize your environment.

1 What is `max_connections` in PostgreSQL?

When managing PostgreSQL at scale or under heavy traffic, understanding connection limits is crucial for performance and stability. One of the most fundamental parameters that controls this is `max_connections`.

Definition

At its core, `max_connections` defines the **maximum number of concurrent client connections** that PostgreSQL will accept at any given time. These connections include:

- **Application users**
- **Database administrators**
- **Replication processes**
- **Background jobs**
- **Monitoring tools**

This setting is essential for controlling **how many active sessions** can interact with the PostgreSQL server simultaneously.

Key Details

- **Default Value:** `100`
This is often sufficient for small to medium-sized applications, but too low for high-concurrency or multi-tenant systems.

- **Includes All Connection Types:**

The count includes **regular users, superusers, background workers**, and **replication roles** — so it's not just about the number of app users.

- **Startup Parameter Only:**

`max_connections` can **only be changed in the `postgresql.conf` file** or via the command line at server startup. You **must restart** the PostgreSQL instance for the change to take effect.

□ **Important Note: More Connections ≠ Better Performance**

It might be tempting to simply raise `max_connections` when facing connection limits—but be cautious. Doing so impacts several key system resources:

- **Memory Usage:**

PostgreSQL reserves a portion of memory for each connection. Increasing the limit too much can exhaust server memory, leading to swapping or crashes.

- **CPU Load:**

More connections often mean more parallel queries, which can saturate CPU cores — especially on OLTP systems.

- **I/O Pressure:**

Simultaneous disk reads/writes from many sessions can strain storage systems, slowing down the entire database.

In many cases, it's better to introduce a **connection pooler** (like **PgBouncer**) instead of drastically increasing `max_connections`.

□ **Best Practices**

- **Monitor:** Use `pg_stat_activity` to check how many connections are active and idle at peak times.

- **Tune Carefully:** Set `max_connections` based on real application demand and available system resources.
- **Use Pooling:** Implement connection pooling for web applications and microservices to reduce idle connections.
- **Scale Horizontally:** If needed, consider scaling out your architecture rather than pushing a single server too far.

By understanding how `max_connections` works and the implications of increasing it, you can ensure your PostgreSQL environment remains stable, performant, and resource-efficient.

2 ☐ Understanding `superuser_reserved_connections` in PostgreSQL

In PostgreSQL, connection management is not just about how many users can access the database — it's also about **ensuring availability for critical operations**. That's where `superuser_reserved_connections` comes in.

This lesser-known but vital configuration parameter plays a key role in maintaining database accessibility, especially during peak load or connection overload scenarios.

☐ What Is `superuser_reserved_connections`?

`superuser_reserved_connections` defines the number of **connection slots reserved exclusively for superusers**—typically database administrators (DBAs).

These reserved slots act as a **safety net**, ensuring that superusers can always access the database, even if all other client connections are maxed out.

✓ Key Characteristics

- **Default Value:** `3`
PostgreSQL sets aside three connections for superusers by default. You can increase or decrease this value in your `postgresql.conf`.
- **Counts Toward** `max_connections`:
These reserved connections are **not in addition** to the `max_connections` setting—they're **included within it**.
- For example:
- If `max_connections = 100`
- And `superuser_reserved_connections = 3`
► Then only **97 connections** are available for regular (non-superuser) clients.
- **Applies Only to Superusers:**
These slots can **only be used by users with superuser privileges**. Regular users will receive a “connection limit exceeded” error if all non-reserved slots are in use.

□ Why It Matters

In production environments, especially those with high traffic, there's always a risk of **hitting the connection ceiling**. When that happens:

- Application clients are blocked.
- Monitoring tools can't connect.
- And **DBAs may be locked out**, unable to even investigate the problem.

Thanks to `superuser_reserved_connections`, **DBAs retain access** even when the system is under pressure. This enables them to:

- Kill idle or runaway sessions
- Investigate blocking or contention

- Perform emergency maintenance or restarts
- Adjust configuration settings to resolve the issue

Without this safeguard, even critical troubleshooting becomes impossible when all connections are exhausted.

□ **Best Practices**

- **Leave the Default or Increase for High-Risk Systems:**
Keep at least 3–5 reserved connections in environments prone to high concurrency or connection spikes.
- **Combine with Connection Pooling:**
Use a tool like PgBouncer to help manage non-superuser connections efficiently and reduce the risk of maxing out connections.
- **Monitor Usage Regularly:**
Query `pg_stat_activity` to understand which sessions are active, and monitor superuser activity separately.

□ **Final Thoughts**

While `max_connections` gets most of the attention, `superuser_reserved_connections` is a **critical safety valve** for PostgreSQL database administrators. By ensuring you never lock yourself out of your own database, this small setting can be the difference between a smooth recovery and a total outage.

3 □ **Why You Should Be Cautious About High `max_connections` in PostgreSQL**

When it comes to PostgreSQL performance tuning, many new administrators are tempted to increase the `max_connections` setting to support more users. After all, more connections must mean more scalability, right?

Not exactly.

While it might seem harmless to allow hundreds — or even thousands — of concurrent database connections, doing so without the right infrastructure and strategy can lead to **severe performance issues**. Let's explore why.

□ **CPU Overhead: The Context Switch Problem**

PostgreSQL follows a **process-per-connection model**. This means every new connection spawns a separate OS-level process. When the number of active sessions grows, so does the overhead on the CPU.

Every process switch requires the operating system to **save the current process state and load another**. This activity is known as a **context switch**.

Why It Matters:

- Context switching increases CPU workload.
- On systems with limited CPU cores, too many concurrent processes can overwhelm the processor.
- The database spends more time **managing processes** than actually executing queries.

This leads to **slow response times**, even when individual queries aren't demanding.

□ **I/O Bottlenecks: Too Many Requests, Too Little Bandwidth**

Each client connection is capable of issuing **read and write operations** — often simultaneously. When hundreds of sessions hit the storage layer concurrently:

- Disk I/O queues grow
- Latency increases

- Cache hit ratios drop
- Checkpoint operations become more expensive

Even high-performance SSDs or SANs can become bottlenecks under heavy concurrent I/O. And for traditional spinning disks, the performance drop can be dramatic.

□ **Session Delays: Resource Contention and Memory Pressure**

More active sessions mean more processes competing for:

- **CPU time**
- **Memory (shared_buffers, work_mem, etc.)**
- **Locks and semaphores**
- **Cache access**

This leads to **performance degradation**, especially when:

- Queries start waiting for locks held by other sessions
- System memory starts swapping
- Parallel query execution stalls due to unavailable worker processes

Essentially, **your system ends up spending more time juggling requests than processing them efficiently.**

□ **Better Alternatives Than Raising `max_connections`**

- ✓ **Use a connection pooler** (e.g., **PgBouncer** or **Pgpool-II**) to manage idle connections more efficiently.
- ✓ **Scale horizontally**: Break workloads across read replicas or use application-level sharding.

- ✓ **Analyze application behavior:** Fix client-side connection leaks or unoptimized usage of long-lived connections.

□ Final Thoughts

Raising `max_connections` seems like a quick fix, but it often masks deeper scalability issues. Without proper planning, high connection counts can **cripple performance** instead of improving it.

Always consider **the impact on CPU, I/O, and memory**, and look toward smarter connection management strategies to maintain a fast, resilient PostgreSQL system.

4 □ How PostgreSQL Allocates Resources for Connections

PostgreSQL is known for its robustness and extensibility — but when it comes to connection handling, **every active session consumes system resources**. Understanding how these resources are allocated is critical for database administrators and DevOps engineers who want to fine-tune performance.

Let's break down how PostgreSQL handles resource allocation for each connection and what trade-offs you need to consider.

□ What Resources Does a Connection Use?

Every **active database connection** in PostgreSQL consumes three key types of memory:

1. □ OS-Level Memory (RAM)

Each connection creates a **dedicated PostgreSQL backend process** at the operating system level. This process consumes memory for the session itself, including stack space and other kernel resources. More connections = more system memory usage.

2. □ PostgreSQL Shared Buffers

Shared buffers (`shared_buffers`) act as a shared memory cache where PostgreSQL stores data blocks read from disk. Every session interacts with this shared memory pool to read/write data. With more connections, **contention increases**, and shared buffers may become a bottleneck unless tuned correctly.

3. ☐ Query-Specific Memory (`work_mem`)

When a query performs operations like sorting, hashing, or joins, it uses per-query memory called `work_mem`. This is **allocated for every operation, per connection**—and potentially multiple times within a single query. If many users are running complex queries simultaneously, total memory consumption can spike quickly.

☐ Rule of Thumb

The more connections you allow, the less memory is available per connection.

This is a crucial design consideration. Simply increasing `max_connections` without adjusting memory parameters can lead to resource exhaustion, poor query performance, or even out-of-memory crashes.

☐ Example Scenario: Tuning Memory Parameters

Let's look at how some key parameters interact when managing connections:

Parameter Impact `max_connections` Higher values mean more concurrent sessions, increasing total memory usage `work_mem` Needs to be carefully tuned if many users run memory-intensive queries `shared_buffers` May need to be increased to accommodate the extra load from more sessions

For example, if you set:

- `max_connections = 500`

- `work_mem = 8MB`

That's potentially **4GB of query memory** in use *per query operation*, not counting overhead or other processes.

If you're not careful, this can easily overwhelm your server's RAM and crash the database.

□ Best Practices

- ✓ **Don't set** `max_connections` **too high** unless absolutely needed. Use connection pooling (e.g., PgBouncer) instead.
- ✓ **Balance** `work_mem` according to expected concurrency and workload complexity.
- ✓ **Monitor shared buffer hit ratios** using tools like `pg_stat_database` or external monitoring platforms.
- ✓ **Always benchmark** memory settings in a staging environment before rolling out to production.

□ Final Thoughts

PostgreSQL gives you fine-grained control over memory and connection behavior — but with great power comes great responsibility. Each connection consumes memory in multiple layers, and a poorly tuned system can become unstable under load.

Instead of throwing more connections at a performance problem, take the time to **understand how memory is allocated and shared** — and tune accordingly.

5 □ Connection Pooling in PostgreSQL: The Secret Weapon for Performance and Stability

One of the most common performance bottlenecks in PostgreSQL isn't the database engine itself — but how applications **connect** to it. Many

developers and DBAs instinctively try to solve connection issues by increasing the `max_connections` setting. But in most cases, there's a far better solution:

✓ **Connection Pooling.**

Let's dive into what it is, why it works, and how to implement it effectively in your PostgreSQL environment.

□ **What Is Connection Pooling?**

A **connection pooler** is a middleware layer that manages a pool of reusable database connections. Instead of opening a new PostgreSQL connection for every client or request, the pooler:

- Maintains a fixed number of open connections to the database
- Reuses them across many application requests
- Reduces the overhead of frequent connection creation and teardown

This helps keep PostgreSQL resource usage low and efficient — even when thousands of users are hitting your application.

✓ **Two Types of Connection Poolers**

1. □ **Application-Level Poolers**

Most modern web frameworks and ORMs (Object-Relational Mappers) offer built-in connection pooling support. These are implemented **within the application layer**, and examples include:

- **SQLAlchemy** (Python)
- **Hibernate** (Java)
- **ActiveRecord** (Rails)

- **Node.js PG-pool** (JavaScript)

These are easy to configure and can often be enabled with just a few lines of code.

2. ☐ **PostgreSQL-Native Poolers**

These tools sit **between your app and the PostgreSQL server**, acting as centralized connection managers. Two popular options:

⚡ **PgBouncer**

- Lightweight and extremely fast
- Ideal for handling large numbers of short-lived connections
- Recommended for most PostgreSQL deployments

☐ **Pgpool-II**

- More feature-rich
- Supports connection pooling, **load balancing**, **failover**, **query caching**, and **replication**
- Better suited for complex, high-availability environments

☐ **The Real Goal of Pooling**

Keep actual PostgreSQL connections low, while allowing your application to scale to many users.

This gives you the best of both worlds:

- Efficient resource usage on the database server
- Scalable performance on the application side

With pooling, you can:

- Serve thousands of concurrent users
- Avoid hitting `max_connections`
- Reduce memory and CPU pressure on PostgreSQL
- Improve app response times during peak loads

□ **Best Practices**

- ✓ Always use PgBouncer or similar for production-grade applications
- ✓ Set pooling mode to `transaction` for best performance in stateless applications
- ✓ Monitor connection usage and fine-tune pool sizes based on traffic patterns
- ✓ Combine with application-level pooling when needed for layered efficiency

□ **Final Thoughts**

Raising `max_connections` might seem like the easy fix—but it’s rarely the right one. Connection pooling is the **smarter, more scalable solution**. Whether you use lightweight tools like **PgBouncer** or feature-rich solutions like **Pgpool-II**, you’ll dramatically improve your PostgreSQL performance, especially under high concurrency.

6 □ **Pre-checks Before Increasing `max_connections` in PostgreSQL**

It’s tempting to resolve “too many connections” errors in PostgreSQL by simply increasing the `max_connections` setting. But doing so without proper preparation can quickly backfire—leading to system instability, crashes, or performance degradation.

Before you touch that number, there are a few **critical pre-checks** you should always perform. Let’s walk through them step by step.

✓1. Ensure Your Application Isn't Leaking Connections

Before increasing `max_connections`, ask yourself: *Does my application actually need more connections—or is it just not closing them properly?*

A **connection leak** happens when the application opens database connections but doesn't close or reuse them. Over time, this saturates the available connections, even if the workload is light.

What to Do:

- Use your application's logs or database monitoring tools to identify long-lived or idle connections.
- Inspect the PostgreSQL `pg_stat_activity` view:

```
SELECT pid, username, application_name, state, backend_start FROM pg_stat_activity;
```

- Fix connection lifecycle management in the application or ORM.

Leaking connections is like leaving water taps running — you don't solve it by opening more pipes.

✓2. Implement Connection Pooling

As discussed in earlier sections, a **connection pooler** (such as PgBouncer) allows you to efficiently reuse a limited number of database connections across many client requests.

Rule of thumb: If you're not using connection pooling, you probably don't need to raise `max_connections`—you need to fix your architecture.

Benefits:

- Reduced resource usage
- Improved query throughput
- Better control over client behavior

Only after implementing pooling should you consider raising `max_connections`, and even then—likely by less than you expected.

✓3. Review OS Kernel Parameters

Raising `max_connections` directly affects how much shared memory PostgreSQL uses. That means the **operating system must be able to support it**.

Key Linux Parameters to Check:

- `shmmax`: Maximum size of a single shared memory segment
- `shmall`: Total amount of shared memory available
- **Semaphores**: PostgreSQL uses semaphores for process coordination

Use the following to inspect:

```
cat /proc/sys/kernel/shmmax
cat /proc/sys/kernel/shmall
ipcs -l
```

If these values are too low, PostgreSQL may fail to start after changing `max_connections`. Always ensure the kernel is properly tuned before applying memory-heavy configuration changes.

✓4. Use Tools Like `pgtune` to Guide Proper Sizing

Rather than guessing, you can use tools like [pgtune](#) to help calculate appropriate settings based on your system's hardware and workload.

Why Use `pgtune`?

- It considers RAM, CPU, and connection counts
- Adjusts related parameters
like `shared_buffers`, `work_mem`, `max_worker_processes`, etc.
- Helps prevent misconfiguration that could overload your system

These tuning tools give you a baseline configuration, which you can refine through monitoring and load testing.

□ Final Thoughts

Increasing `max_connections` isn't just a one-line change in `postgresql.conf`—it's a **system-level decision** with far-reaching implications. Before raising the number, you must:

- Fix the real root causes (like connection leaks)
- Implement efficient connection pooling
- Ensure your OS can handle the extra load
- Rely on intelligent tuning, not guesswork

Only then can you safely scale your PostgreSQL deployment to meet growing demand — without compromising performance or stability.

7 □ How to Safely Increase `max_connections` in PostgreSQL

Raising the `max_connections` setting in PostgreSQL can help support more concurrent users—but it's not just a matter of changing a number. PostgreSQL relies on shared memory, and your operating system must be properly configured to support the extra load.

In this guide, we'll walk through the **safe and correct way to increase `max_connections`**, including kernel and PostgreSQL configuration steps.

✓Step 1: Verify Current Kernel Shared Memory

PostgreSQL uses shared memory (`shared_buffers`) to cache data pages. When you increase `max_connections`, you may also need to increase `shared_buffers`, which in turn depends on your operating system's **shared memory limit** (`shmmax`).

To check your current `shmmax` value:

```
cat /proc/sys/kernel/shmmax
```

This outputs the maximum size (in bytes) of a single shared memory segment the Linux kernel will allow. If this value is too low, PostgreSQL may fail to start after increasing memory settings.

✓Step 2: Update Kernel Settings (if required)

If your current `shmmax` is insufficient (e.g., below 1–2 GB), you need to increase it.

Temporary Change (until reboot):

```
sudo sysctl -w kernel.shmmax=10066329600
```

This sets the maximum shared memory segment size to 10 GB.

Persistent Change (survives reboots):

```
echo "kernel.shmmax=10066329600" | sudo tee -a /etc/sysctl.conf  
sudo sysctl -p
```

You can also adjust other memory parameters here, such as `shmall`, depending on your system usage.

✔ Step 3: Modify PostgreSQL Parameters

Now that your OS supports the required memory, it's time to update PostgreSQL's configuration file:

Open `postgresql.conf` (typically located in `/etc/postgresql/<version>/main/` or `$PGDATA`):

```
sudo nano /etc/postgresql/14/main/postgresql.conf
```

Update the following parameters:

```
max_connections = 200
shared_buffers = 256MB # Adjust based on total system RAM
```

❑ **Tip:** For high `max_connections` values, consider increasing `work_mem`, `max_locks_per_transaction`, and other memory-related settings too.

✔ Step 4: Restart PostgreSQL

For the changes to take effect, restart the PostgreSQL service:

```
sudo systemctl restart postgresql
```

You can confirm the new settings by connecting to the database and running:

```
SHOW max_connections;  
SHOW shared_buffers;
```

□ Final Thoughts

Increasing `max_connections` in PostgreSQL is not just a database tweak—it's a **system-wide memory adjustment** that must be done carefully. Following the steps above ensures:

- Your operating system can handle the memory PostgreSQL needs
- The database is tuned for stability
- You avoid crashes, startup failures, or degraded performance

Always test these changes in a staging environment before applying them in production, and combine this with connection pooling for best results.

8 □ Monitoring Connections in PostgreSQL: Stay Ahead of Limits

One of the simplest yet most powerful ways to keep your PostgreSQL database running smoothly is by **actively monitoring connection usage**. Running out of available connections can lead to application errors, locked-out DBAs, and stalled services — especially if you're unaware that you're nearing your limits.

In this post, we'll go over essential SQL queries that help you **track, analyze, and prevent connection-related bottlenecks**.

□ Why Monitor Connections?

PostgreSQL enforces a strict limit on the number of simultaneous client connections via the `max_connections` setting. Exceeding this limit causes new connection attempts to fail with errors like:

```
FATAL: remaining connection slots are reserved for non-
replication superuser connections
```

To prevent this, you need visibility into how many connections are in use — **and how close you are to the ceiling.**

✓ Essential Queries for Connection Monitoring

Here are the three key metrics every PostgreSQL admin should monitor regularly:

□ 1. Total Active Connections

To get the current number of connections (both idle and active):

```
SELECT count(*) AS active_connections
FROM pg_stat_activity;
```

This query gives you a snapshot of how many sessions are currently connected to your database. If this number frequently approaches your `max_connections`, it's time to investigate further.

□ 2. Superuser Reserved Slots

PostgreSQL reserves a few connection slots specifically for superusers, which are critical for emergency access.

Check how many are reserved:

```
SELECT setting::int AS superuser_reserved
FROM pg_settings
WHERE name = 'superuser_reserved_connections';
```

Note: These slots are included within the `max_connections` limit. So, if `max_connections = 100` and `superuser_reserved_connections = 3`, only **97** slots are available to regular users.

□ 3. Maximum Allowed Connections

To confirm your configured connection ceiling:

```
SELECT setting::int AS max connections
FROM pg_settings
WHERE name = 'max_connections';
```

Knowing this value allows you to gauge your connection usage percentage:

```
-- Optional: Quick usage percentage calculation
SELECT
  count(*) * 100.0 / (SELECT setting::int FROM pg_settings WHERE name =
    'max_connections')
  AS connection usage percent
FROM pg_stat_activity;
```

□ Best Practices

- ✓ **Automate alerts** using monitoring tools like **Prometheus + Grafana**, **pgMonitor**, or **Zabbix**.
- ✓ **Set thresholds** (e.g., alert if usage > 80%).
- ✓ **Use connection pooling** to reduce idle connections.
- ✓ **Investigate idle or long-running sessions** that may unnecessarily hold connections.

□ Final Thoughts

Proactive connection monitoring helps you avoid one of the most frustrating database issues: **hitting connection limits at the worst possible time**. By regularly running a few simple queries — or

integrating them into your monitoring stack — you gain visibility and control over one of PostgreSQL’s most important performance levers.