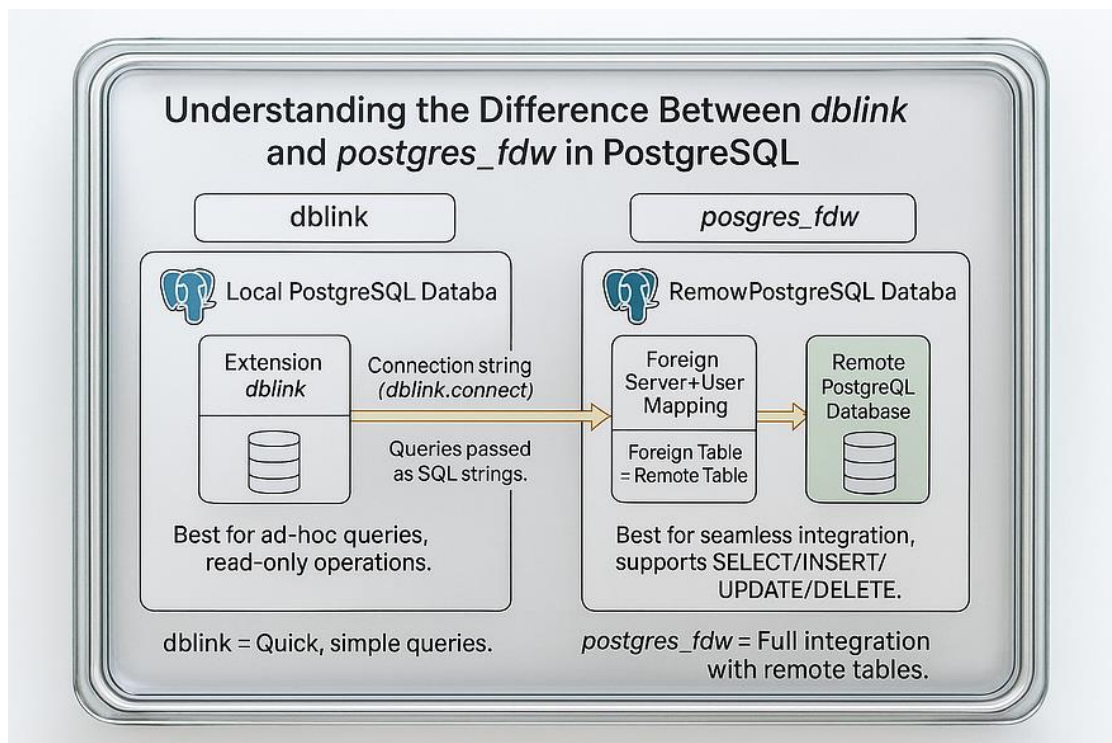


Understanding the Difference Between `dblink` and `postgres_fdw` in PostgreSQL



PostgreSQL is not just a relational database; it is a highly extensible system. Two of its most commonly used extensions for working across multiple PostgreSQL databases are `dblink` and `postgres_fdw` (**Foreign Data Wrapper**).

At first glance, they might look similar because both allow you to connect to another PostgreSQL database. But in reality, they serve different purposes and have different strengths.

□ 1. What is `dblink`?

The `dblink` **extension** is one of the oldest ways to connect one PostgreSQL database to another. It allows you to run **queries as text strings** against a remote database.

- Think of it as a **bridge**: you connect to the other database, send a query as text, and get the results back.
- You must manually define the expected column structure when pulling data.
- It's lightweight but also limited — best suited for **ad-hoc queries**.

Example workflow with `dblink`:

1. Install and enable the extension (`CREATE EXTENSION dblink`).
2. Configure a remote connection.
3. Run a query like:

```
SELECT *
FROM dblink('conn db link', 'SELECT item id, quantity, price FROM items')
AS result(item_id INT, quantity INT, price NUMERIC(10,2));
```



```
postgres=# SELECT *
postgres=# FROM dblink('conn db link', 'SELECT item id, quantity, price FROM items')
postgres=# AS result(item_id INT, quantity INT, price NUMERIC(10,2));
 item id | quantity | price
-----+-----+-----
 9093121 |      297 | 395.80
 9093122 |      640 | 226.10
 9093123 |      238 | 173.32
 9093124 |      200 |  80.91
 9093125 |       18 | 222.13
 9093126 |      945 | 101.12
 9093127 |      419 | 192.06
 9093128 |      739 | 493.85
 9093129 |      755 | 163.90
 9093130 |      374 |  69.76
 9093131 |      613 | 339.56
 9093132 |      923 |  75.71
 9093133 |      917 | 461.97
 9093134 |      916 | 156.12
 9093135 |      719 | 433.66
```

Here, the SQL query is written as a string and executed on the remote server.

□ 2. What is `postgres_fdw`?

The `postgres_fdw` **extension** is a modern, more powerful replacement for many use cases of `dblink`. FDW stands for **Foreign Data Wrapper**, which is part of PostgreSQL's implementation of the SQL/MED standard (Management of External Data).

- Instead of sending queries as strings, you can create **foreign tables** in your local database that directly map to remote tables.
- These foreign tables look and behave like normal local tables.
- You can **SELECT, INSERT, UPDATE, and DELETE** records across databases.
- PostgreSQL's **query planner** can optimize queries across local and remote data sources.

Example workflow with `postgres_fdw`:

1. Install and enable the extension (`CREATE EXTENSION postgres_fdw`).
2. Create a foreign server that points to the remote database.
3. Create a user mapping with credentials.
4. Define a foreign table that mirrors the remote table.
5. Run queries against it directly, without writing raw SQL strings.

```
SELECT * FROM items_remote WHERE price > 100;
```

The query above looks like a local query, but it is executed on the remote database behind the scenes.

□ 3. Key Differences at a Glance

Feature	<code>dblink</code>	<code>postgres_fdw</code>	Query style	SQL passed as a string	Normal SQL on foreign tables
Integration level	Lightweight, ad-hoc	Full integration with local schema	Operations supported	Mostly SELECT (writes are clunky)	SELECT, INSERT, UPDATE, DELETE
Query optimization	No	Yes (Postgres planner pushes conditions to remote server)	Ease of use	Simple for quick queries	More setup, but easier for ongoing use
Use case	One-time lookups, quick joins	Long-term integration, production workloads			

□ In short:

- Use `dblink` if you just need a quick lookup across databases.
- Use `postgres_fdw` if you need continuous, reliable, and optimized access to remote tables.

Step-by-Step Guide: Implementing `postgres_fdw` in PostgreSQL 15

Now that we understand the difference, let's dive into the **step-by-step process** of setting up and using `postgres_fdw` in PostgreSQL 15.

Step 1: Install the contrib package

If you haven't already installed contrib modules for PostgreSQL 15, do so:

```
sudo yum install postgresql15-contrib
```

```
[root@ggnode1 dbs]# sudo yum install postgresql15-contrib    # or: sudo yum install
postgresql15-contrib
o17_UK6R6
| 3.0 kB  00:00:00
```

```

ol7 addons
| 3.0 kB 00:00:00
ol7 developer
| 3.0 kB 00:00:00
ol7_developer_EPEL
| 3.6 kB 00:00:00
ol7_latest
| 3.6 kB 00:00:00
ol7_optional_latest
| 3.0 kB 00:00:00
ol7_software_collections
| 3.0 kB 00:00:00
pgdg-common/7Server/x86_64/signature
| 665 B 00:00:00
pgdg-common/7Server/x86_64/signature
| 2.9 kB 00:00:00 !!!
pgdg12/7Server/x86_64/signature
| 665 B 00:00:00
pgdg12/7Server/x86_64/signature
| 3.6 kB 00:00:00 !!!
pgdg13/7Server/x86_64/signature
| 665 B 00:00:00
pgdg13/7Server/x86_64/signature
| 3.6 kB 00:00:00 !!!
pgdg14/7Server/x86_64/signature
| 665 B 00:00:00
pgdg14/7Server/x86_64/signature
| 3.6 kB 00:00:00 !!!
pgdg15/7Server/x86_64/signature
| 665 B 00:00:00
pgdg15/7Server/x86_64/signature
| 3.6 kB 00:00:00 !!!
Resolving Dependencies
--> Running transaction check
--> Package postgresql15-contrib.x86_64 0:15.14-1PGDG.rhel7 will be installed
--> Processing Dependency: libpython3.6m.so.1.0()(64bit) for package: postgresql15-
contrib-15.14-1PGDG.rhel7.x86_64
--> Running transaction check
--> Package python3-libs.x86_64 0:3.6.8-21.0.1.el7_9 will be installed
--> Processing Dependency: python(abn) = 3.6 for package: python3-libs-3.6.8-
21.0.1.el7_9.x86_64
--> Running transaction check
--> Package python3.x86_64 0:3.6.8-21.0.1.el7_9 will be installed
--> Processing Dependency: python3-pip for package: python3-3.6.8-
21.0.1.el7_9.x86_64
--> Processing Dependency: python3-setuptools for package: python3-3.6.8-
21.0.1.el7_9.x86_64
--> Running transaction check
--> Package python3-pip.noarch 0:9.0.3-8.0.3.el7 will be installed
--> Package python3-setuptools.noarch 0:39.2.0-10.0.3.el7 will be installed
--> Finished Dependency Resolution

```

Dependencies Resolved

```

=====
=====
=====
Package
Version
Size
Arch
Repository
=====
=====
Installing:
  postgresql15-contrib                x86_64
  15.14-1PGDG.rhel7                  pgdg15
  722 k
Installing for dependencies:
  python3                            x86_64
  3.6.8-21.0.1.el7_9                 ol7_latest
  70 k
  python3-libs                        x86_64
  3.6.8-21.0.1.el7_9                 ol7_latest
  7.0 M

```

```
python3-pip                                noarch
9.0.3-8.0.3.el7                            ol7_latest
1.6 M
python3-setuptools                        noarch
39.2.0-10.0.3.el7                         ol7_latest
629 k
```

Transaction Summary

```
=====
=====
=====
Install 1 Package (+4 Dependent packages)
```

Total download size: 10 M

Installed size: 50 M

Is this ok [y/d/N]: y

Downloading packages:

```
(1/5): python3-3.6.8-21.0.1.el7_9.x86_64.rpm
| 70 kB 00:00:00
(2/5): postgresql15-contrib-15.14-1PGDG.rhel7.x86_64.rpm
| 722 kB 00:00:00
(3/5): python3-pip-9.0.3-8.0.3.el7.noarch.rpm
| 1.6 MB 00:00:00
(4/5): python3-setuptools-39.2.0-10.0.3.el7.noarch.rpm
| 629 kB 00:00:00
(5/5): python3-libs-3.6.8-21.0.1.el7_9.x86_64.rpm
| 7.0 MB 00:00:01
```


Total

This package contains both `dblink` and `postgres_fdw`.

Step 2: Enable the extension in your database

Log in to PostgreSQL and enable the FDW extension:

```
CREATE EXTENSION postgres_fdw;
```

```
postgres=# CREATE EXTENSION postgres_fdw;
CREATE EXTENSION
postgres=#
```

Verify that it's installed:

```
\dx postgres_fdw
```

```

postgres=# \dx postgres fdw
               List of installed extensions
  Name          | Version | Schema | Description
-----+-----+-----+-----
--
postgres_fdw | 1.1      | public | foreign-data wrapper for remote PostgreSQL
servers
(1 row)

postgres=#

```

Step 3: Create a foreign server

Define a connection to the remote PostgreSQL server:

```

CREATE SERVER server_remote FOREIGN DATA WRAPPER postgres_fdw
postgres=# OPTIONS (host '192.168.50.11', dbname 'postgres', port '5432');

postgres=# CREATE SERVER server_remote FOREIGN DATA WRAPPER postgres_fdw
postgres=# OPTIONS (host '192.168.50.11', dbname 'postgres', port '5432');
CREATE SERVER
postgres=#

```

Here:

- **host** = IP address of the remote server
- **dbname** = remote database name
- **port** = default PostgreSQL port (5432)

Step 4: Create a user mapping

Map the current user to the remote user by providing login credentials:

```

CREATE USER MAPPING FOR current_user SERVER server_remote
OPTIONS (user 'postgres', password 'oracle123');

```

```
postgres=# CREATE USER MAPPING FOR current user SERVER server remote
postgres=# OPTIONS (user 'postgres', password 'oracle123');
CREATE USER MAPPING
postgres=#
```

This ensures that when the local user connects to the foreign server, it uses the specified username and password.

Step 5: Import schema or create foreign tables

At this point, you can either **import an entire schema** from the remote database or manually define a foreign table.

Option A: Import all tables from a schema

```
IMPORT FOREIGN SCHEMA public FROM SERVER server_remote
INTO public;

postgres=# IMPORT FOREIGN SCHEMA public FROM SERVER server_remote
postgres=# INTO public;
IMPORT FOREIGN SCHEMA
postgres=#
```

This automatically creates local foreign tables for everything in the remote `public` schema.

Viewing Local and Foreign Tables

After setting up `postgres_fdw` and importing or creating foreign tables, you can list all relations (tables, views, sequences, etc.) in your current schema. In PostgreSQL, the command for this is:

```
\d+
```


Example output:

```
postgres=# \dt+

```

Schema	Name	Type	List of relations		
method	Size	Description	Owner	Persistence	Access
public	items	foreign table	postgres	permanent	
0 bytes					
public	orders	foreign table	postgres	permanent	
0 bytes					
public	products	table	postgres	permanent	heap
977 MB					
public	products product id seq	sequence	postgres	permanent	
8192 bytes					
(4 rows)					

```
postgres=#
```

What This Means

- **items (foreign table):**
This is a table that resides on the remote PostgreSQL server but is accessible locally through `postgres_fdw`.
- **orders (foreign table):**
Another remote table, seamlessly integrated into the local database. Queries on this table are executed against the remote database.
- **products (local table):**
A regular table that exists locally in the `public` schema.
- **products_product_id_seq (sequence):**
A sequence object used to generate IDs for the `products` table.

Key Takeaway

With `postgres_fdw`, your local PostgreSQL instance can contain a **mix of local and remote objects**.

- To your application, these look almost the same.
- The big difference is that **foreign tables delegate operations to the remote database**.

This setup makes it possible to query and join data across multiple databases without complex ETL processes.

Option B: Define a specific foreign table

If you want to add new tables after importing a foreign schema, you cannot re-import the foreign tables. Instead, you can use this option to add additional tables.

```
postgres=# IMPORT FOREIGN SCHEMA public FROM SERVER server_remote
postgres=# INTO public;
ERROR:  relation "items" already exists
CONTEXT:  importing foreign table "items"
postgres=#
```

```
CREATE FOREIGN TABLE items_remote (
    item_id INT,
    quantity INT,
    price NUMERIC(10,2)
)
SERVER server_remote
OPTIONS (table_name 'items');
```

```
postgres=# CREATE FOREIGN TABLE items_remote (
postgres=#     item_id INT,
postgres=#     quantity INT,
postgres=#     price NUMERIC(10,2)
postgres=# )
postgres=# SERVER server_remote
postgres=# OPTIONS (table_name 'items');
CREATE FOREIGN TABLE
postgres=#
```

Here, `items_remote` is the local representation of the remote `items` table or any new tables.

```
postgres=# \dt

```

Schema	Name	Type	Owner	Persistence	Access
method	Size	Description			
public	items	foreign table	postgres	permanent	
public	items_remote	foreign table	postgres	permanent	

```

public | orders | foreign table | postgres | permanent |
| 0 bytes |
public | products | table | postgres | permanent | heap
| 977 MB |
public | products_product_id_seq | sequence | postgres | permanent |
| 8192 bytes |
(5 rows)

postgres=#

```

Step 6: Query the foreign table

Now you can query the remote table as if it were local:

```

SELECT * FROM items WHERE price > 200;

postgres=# SELECT * FROM items WHERE price > 200;

 item_id | quantity | price
-----+-----+-----
 9093121 |      297 | 395.80
 9093122 |      640 | 226.10
 9093125 |       18 | 222.13
 9093128 |      739 | 493.85
 9093131 |      613 | 339.56
 9093133 |      917 | 461.97
 9093135 |      719 | 433.66
 9093137 |      211 | 212.29
 9093138 |      732 | 212.20
 9093140 |      319 | 295.23

```

Step 7: Perform DML operations

Unlike `dblink`, `postgres_fdw` allows you to modify data on the remote server.

- **Delete data:**

```

DELETE FROM items WHERE item_id = 9093129;

postgres=# DELETE FROM items WHERE item_id = 9093129;
DELETE 1
postgres=#

```

- **Insert data:**

```
INSERT INTO items VALUES (9093129, 500, 250.00);
```

```
postgres=# INSERT INTO items VALUES (9093129, 500, 250.00);  
INSERT 0 1  
postgres=#
```

- **Update data:**

```
UPDATE items SET price = 199.99 WHERE item_id = 9093121;
```

```
postgres=# UPDATE items SET price = 199.99 WHERE item_id = 9093121;  
UPDATE 1  
postgres=#
```

All changes are applied directly to the remote database.

☐ Practical Considerations

1. Performance:

- `postgres_fdw` pushes down WHERE conditions to the remote server, reducing network traffic.
- Joins between local and foreign tables may be less efficient, so use filters.

2. Security:

- Avoid storing plain-text passwords in user mappings. Use role-based authentication or password files.

3. Use cases:

- Ideal for distributed reporting, microservices accessing shared data, or consolidating analytics.
- Not a replacement for replication — use logical or streaming replication for high availability.

□ Final Thoughts

- `dblink` is the older, simpler way: best for **occasional queries**.
- `postgres_fdw` is the modern, integrated way: best for **production systems where remote tables should behave like local ones**.

With the steps above, you can seamlessly connect PostgreSQL databases using `postgres_fdw` and perform both read and write operations across multiple systems.