

Logical Replication final

Publication and Subscription:-

- **Key Concepts**

- **Publisher:-**

The database instance that makes data available for replication. It defines what data should be replicated through publications.

- **Subscriber:-**

The database instance that receives replicated data from publishers by subscribing to publications.

- **Publication:-**

A set of changes generated from a table or a group of tables, optionally filtered by specific operations (INSERT, UPDATE, DELETE).

- **Subscription:-**

The downstream side of logical replication. It defines the connection to another database and set of publications to subscribe to.

1. check logical replication status by using-

```
Select * from pg_replication_stats where active=true;
```

Configuration file Publisher(primary)

Critical `postgresql.conf` Settings for Publisher

Enable logical replication
`wal_level = logical`

Number of concurrent WAL sender processes (each subscriber uses 1)
`max_wal_senders = 400`

Number of replication slots allowed (logical and physical)
`max_replication_slots = 400`

Number of logical replication workers for subscriptions (per database)
`max_logical_replication_workers = 300` # Pick ONE value, don't duplicate

Number of background workers for logical replication
`max_sync_workers_per_subscription = 150` # Optional but useful for parallel table sync

Total number of background worker processes
`max_worker_processes = 800`

Total number of parallel workers
`max_parallel_workers = 400`

- **Create user:-**

```
CREATE USER replicator WITH REPLICATION ENCRYPTED PASSWORD 'your_very_strong_password';
```

- **Create publication on the primary server:**

```
CREATE PUBLICATION pub_<schema_name> FOR TABLE <schema_name.table_name>,  
<schema_name.table_name>, <schema_name.table_name> WITH (publish = 'insert, update, delete, truncate',  
publish_via_partition_root = false);
```

- **Publications define what data to replicate:-**

-- Create a publication for all tables

```
CREATE PUBLICATION all_tables FOR ALL TABLES;
```

-- Or create selective publications

```
CREATE PUBLICATION products_only FOR TABLE products;
```

```
CREATE PUBLICATION orders_only FOR TABLE orders;
```

-- Create a publication with specific operations

```
CREATE PUBLICATION products_inserts_only FOR TABLE products WITH (publish = 'insert');
```

-- View existing publications

```
SELECT * FROM pg_publication;-- View tables in publications
```

```
SELECT * FROM pg_publication_tables;
```

- **Grant permissions to the replication user on the primary server**

```
GRANT CONNECT ON DATABASE <database_name> TO replicator;
```

```
GRANT USAGE ON SCHEMA <schema_name> to replicator;
```

```
GRANT ALL ON ALL TABLES IN SCHEMA <schema_name> TO replicator;
```

```
GRANT ALL ON ALL SEQUENCES in schema <schema_name> TO replicator;
```

- **wal_level = logical:** Absolutely essential. This enables the generation of logical decoding information in the WAL, which is what logical replication uses.
- **max_worker_processes:** Set high enough to accommodate all background workers, including replication. A good rule of thumb: `max_worker_processes >= max_logical_replication_workers (on subscriber) + 1` (or more, depending on other background tasks).
- **max_wal_senders:** The maximum number of concurrent connections from standby servers or logical replication subscribers. It should be greater than or equal to `max_replication_slots` plus any physical replicas you might have.
- **max_replication_slots:** The maximum number of replication slots that can be active at the same time. This should be at least equal to the number of subscriptions you plan to have, plus any additional slots for table synchronization workers if you divide publications.

Note:-Important constraints:

- `max_logical_replication_workers` should not exceed `max_worker_processes`.
- `max_sync_workers_per_subscription` should be less than or equal to `max_logical_replication_workers`.

Explanation:-

❖ System Resources (CPU, RAM, I/O):-

- Each replication slot, WAL sender, and logical worker consumes CPU and memory.
- Don't over-allocate these values beyond what your system can handle.
- Rule of thumb: For high values (e.g., 400), ensure the server has enough CPUs (ideally 1 core per 10–20 replication workers) and ample RAM.

❖ `wal_level = logical`

- Enables logical decoding of WAL (Write-Ahead Log).
- Required for publishing/subscribing changes to tables.
- Slightly increases WAL size compared to replica, due to additional information needed for decoding.

❖ `max_wal_senders`

- Number of WAL sender processes allowed.
- Each replication connection (logical or physical) uses one.
- Needed for both streaming replication and logical replication.
- Too many unused slots can cause **WAL bloat** (old WAL files won't be removed until acknowledged).

❖ `max_replication_slots`

- Number of replication slots PostgreSQL will allow.
- Logical replication uses a slot for each subscription to track WAL position.
- One **WAL sender process** is used per active subscription.

❖ `max_logical_replication_workers`

- Total number of logical replication workers (background workers handling subscriptions).

- Higher value allows more concurrent logical subscriptions.
- `max_sync_workers_per_subscription` (optional but useful)
- Number of tables synchronized in parallel during initial sync.
- Speeds up initial table copy when a subscription is created.
- **`max_logical_replication_workers`**: total workers across all DBs.
- **`max_sync_workers_per_subscription`**: controls how many tables are synced in parallel during initial subscription.
- Higher values = faster sync, but more load on the system

❖ **`max_worker_processes`**

- Total background workers allowed by PostgreSQL.
- Logical replication workers, parallel workers, and extensions all count toward this.

❖ **`max_parallel_workers`**

- Maximum parallel workers that can be used by queries.
- Not directly tied to logical replication but important for overall performance.

Configuration file Subscriber

Critical postgresql.conf Settings for Subscriber

Enable logical replication

wal_level = logical

Number of concurrent WAL sender processes (each subscriber uses 1)

max_wal_senders = 800

Number of replication slots allowed (logical and physical)

max_replication_slots = 800

Number of logical replication workers for subscriptions (per database)

max_logical_replication_workers = 800 # Pick ONE value, don't duplicate

Number of background workers for logical replication

max_sync_workers_per_subscription = 200 # Optional but useful for parallel table sync

Total number of background worker processes

max_worker_processes = 1600

Total number of parallel workers

max_parallel_workers = 400

Step 1: Create the Exact Database on Subscriber

```
CREATE DATABASE parkdepot; -- Use your actual database name
```

Step 2: Extract Global Objects (Roles, Tablespaces) from Publisher

```
pg_dumpall -h <publisher_endpoint> -p 5432 -U <admin_user> --globals-only > "globals-roles-tablespaces.sql"
```

Step 3: Extract Only the Schema (Structure) from Publisher

#Schema backup

Options:

-s=schema, -O=without owner, -x=without privileges, -v=verbose

```
pg_dump -h 172.31.27.179 -U postgres -d sourcedb -s -O -x -p 5432 -v > source_db_schema.sql
```

```
pg_dump -h <publisher_IP> -U <admin_user> -d <database_name> --schema-only > schema.sql
```

Step 5: Import Global Objects and Schema to Subscriber

```
psql -U <admin_user> -h <subscriber_IP> -p 5432 -d postgres -f globals-roles-tablespaces.sql
```

```
psql -h <subscriber_IP> -U <admin_user> -d <database_name> -f schema.sql
```

Step 4: Create the Subscription

- **Create subscription on the secondary server:**

```
CREATE SUBSCRIPTION sub_<schema_name>
```

```
    CONNECTION 'host=<host_ip_address> port=<host_port> user=replicator
dbname=<database_name> connect_timeout=10 password=<password>
sslmode=prefer'
```

```
    PUBLICATION pub_<schema_name>
```

```
    WITH (connect = true, enabled = true, copy_data = true, create_slot = true,
synchronous_commit = 'off');
```

- **Grant permissions to the secondary server:**

```
GRANT CONNECT ON DATABASE ecgc_reporting_db TO rep_user, smile_opr;
```

```
GRANT USAGE ON SCHEMA <schema_name> TO rep_user, smile_opr; GRANT
```

```
SELECT ON ALL TABLES IN SCHEMA <schema_name> TO rep_user;
```

-- View subscription status

```
SELECT * FROM pg_subscription;
```

-- Check subscription worker status

```
SELECT * FROM pg_stat_subscription;
```

-- View replication slots on the publisher-- (Run this on the publisher)

```
SELECT * FROM pg_replication_slots;
```

Monitoring Your Logical Replication: The Lifeline

- **On the Publisher (pg_stat_replication_slots):**

```
SELECT slot_name, active, wal_status, wal_delay_lags_bytes, restart_lsn, confirmed_flush_lsn FROM pg_stat_replication_slots;
```

- **On the Subscriber (pg_stat_subscription and pg_stat_subscription_workers):**

```
SELECT subname, subenabled, subconninfo, substate, sublag FROM pg_stat_subscription; SELECT subid, relid, last_applied_lsn, last_received_lsn, wal_records, wal_bytes, sync_state FROM pg_stat_subscription_workers;
```

Logical Replication command:-

- **create ,alter:- (DDL command)**

1. first we have to run on reporting db

2. run on production db

3.add this table on publication

4. refresh subscription:- (You **must run this on the Subscriber database**, not on the Publisher.)

```
alter subscription <subscription_name> refresh publication
```


=====

● **publisher:- to see lag in GB**

```
SELECT
slot_name,
active,
restart_lsn,
pg_size_pretty(pg_wal_lsn_diff(pg_current_wal_lsn(), restart_lsn)) AS wal_lag_pretty,
round(pg_wal_lsn_diff(pg_current_wal_lsn(), restart_lsn) / 1024.0 / 1024.0, 2) AS wal_lag_mb,
round(pg_wal_lsn_diff(pg_current_wal_lsn(), restart_lsn) / 1024.0 / 1024.0 / 1024.0, 2) AS
wal_lag_gb
FROM pg_replication_slots
WHERE slot_type = 'logical' AND active = false;
```

subscriber:- to see lag

```
SELECT
slot_name AS replication_slot,
active AS is_active,
restart_lsn,
pg_size_pretty(pg_wal_lsn_diff(pg_current_wal_lsn(), restart_lsn)) AS wal_lag_pretty,
ROUND(pg_wal_lsn_diff(pg_current_wal_lsn(), restart_lsn) / 1024.0 / 1024.0, 2) AS
wal_lag_mb,
ROUND(pg_wal_lsn_diff(pg_current_wal_lsn(), restart_lsn) / 1024.0 / 1024.0 / 1024.0, 2) AS
wal_lag_gb
FROM pg_replication_slots
WHERE slot_type = 'logical'
AND active = false;
```

```
SELECT
slot_name AS replication_slot,
active AS is_active,
restart_lsn,
pg_size_pretty(pg_wal_lsn_diff(pg_current_wal_lsn(), restart_lsn)) AS wal_lag_pretty,
ROUND(pg_wal_lsn_diff(pg_current_wal_lsn(), restart_lsn) / 1024.0 / 1024.0, 2) AS
wal_lag_mb,
ROUND(pg_wal_lsn_diff(pg_current_wal_lsn(), restart_lsn) / 1024.0 / 1024.0 / 1024.0, 2) AS
wal_lag_gb
FROM pg_replication_slots
WHERE slot_type = 'logical'
AND active = false;
```

```

SELECT
    subname,
    pid,
    received_lsn,
    latest_end_lsn,
    pg_wal_lsn_diff(latest_end_lsn, received_lsn) AS lag_bytes
FROM pg_stat_subscription
WHERE subname = 'sub_buyer';

```

```

CREATE SUBSCRIPTION sub_intl_bank
CONNECTION 'host=172.16.34.25 port=5432 user=replicator password=prod#replicator
dbname=bud_intl_bank_db connect_timeout=10 sslmode=prefer'
PUBLICATION pub_intl_bank
WITH (
    connect = true,
    enabled = true,
    slot_name = 'sub_intl_bank',
    create_slot = false,
    copy_data = false,
    synchronous_commit = 'off'
);

```

=====

- **To see repliation slot is failed**

```
select * from pg_replication_slots where active=false;
```

- **To refreash subscription:-**

Alter subscription sub_issue_endorse refresh publication;

- **To see lock :-**

```

SELECT
    t.relname AS locked_table,
    l.transactionid,
    l.mode,
    l.pid,
    l.granted
FROM

```

```
pg_catalog.pg_locks l
JOIN pg_catalog.pg_class t ON t.oid = l.relation
WHERE
    t.relname = 'ecgc_ecib_claim_recovery_int_seq_no_tbl';-----How to check table is locked or
unlocked
```

- **To see database size**

```
SELECT
    d.datname AS database_name,
    pg_size_pretty(pg_database_size(d.datname)) AS size
FROM
    pg_database d;--SQL query to get the sizes of all databases
```

```
SELECT pg_size_pretty(SUM(pg_database_size(datname))) AS total_size
FROM pg_database;--Total size of database
```